

Contents

1	README	2
2	Makefile	4
3	osm.cpp	5

1 README

```
1  ransha
2  Ran Shaham (203781000)
3  EX: 1
4
5  FILES:
6  README -- This file
7  osm.cpp -- a file with some code
8  Makefile -- a makefile that creates the library (make with no arguments, or
9              'all')
10
11  NOTES -
12  There are A LOT of magic numbers in my code, I'm hoping that you'll believe me
13  when I say this is for the sake of accuracy in time measurements and not
14  because of laziness or bad coding practice.
15
16  TASK 1:
17  The program takes an argument. It then creates a directory called 'Welcome',
18  inside it it creates another directory called 'To', and inside it creates a
19  file name OS with write attributes (open O_CREAT| O_WRONLY).
20  It then writes some text to this file
21  (prompts us to read the course guidelines).
22  I also noticed that the number of bytes written to the file is dependant by
23  the length of the arguments I supplied.
24  Namely, with 1 letter argument the number of bytes written is 178 and each
25  added letter adds another byte. Then, the program closes and deletes the file
26  and remove both directories.
27
28  I learned it by 'man'ing the sys calls:
29  mkdir - creates a directory, returns 0 upon success
30  open - opens a file with the attributes set by the flags give
31  file descriptor
32  write - writes to a file given by its file descriptor from a given pointer a
33  given number of
34  bytes.
35  close - closes an open file
36  unlink - deletes a file (deletes a name from the filesystem?) given by its
37  descriptor.
38  rmdir - removes an empty directory
39
40  TASK 2:
41  osm_operation_time - in this function I used loop unrolling to perform a lot
42  of basic operations on random numbers (they were randomized before the
43  measurement start).
44  I then divided by the number of total operations performed, hoping that the
45  loop operations weren't significant compared to the operations.
46  Note that I didn't assign the result of the operations to a variable so I
47  won't count the assignment time, I'm aware this could be problematic - if the
48  compiler decides that an unassigned rhs expression should be ignored, couldn't
49  think of a better way.
50  The loop unrolling, as suggested in the ex description, is to get average time
51  bigger than 0 because basic operations are performed quickly.
52
53  osm_function_time - I created two empty functions in my cpp file and called
54  them one by one a lot of times in each iteration (~500 iterations in each
55  main for iteration, in each one of those I made 12 function calls).
56  Again, this is to avoid manipulating very small numbers, so i'd get more
57  accurate results.
58
59  osm_syscall_time - in each iteration I made 5 empty system calls. This is
```

-3/-15 Wrong answer(s) in the README.
(code='README_general_err') Partial
explanation of the system calls including
parameters and return value

60 much fewer than the ones above because these functions take much longer time
61 so the measured numbers will be reasonable to handle.
62
63 osm_disk_time - Most of my time was spent HERE. I tried to write and read from
64 files that I opened with O_SYNC and O_DIRECT flags to avoid caching and
65 buffering. It was very hard since the writing and reading failed - after
66 reading a lot about it I figured it's because of alignment issues, couldn't
67 understand it, not for a bit.
68 So I removed the flags and measured the times without them. I'm aware this
69 measurement isn't the best, but it got reasonable results (~900 ns which is
70 much higher than all the others) so I left it at that.
71 I used two files hoping that accessing them alternatively will prevent the
72 disk from optimizing the access to their locations.
73
74 Note that the files I handled were created and deleted by the osm_init,
75 osm_finalizer functions.

2 Makefile

You may use built-in variable CXX which defaults to g++.

```
1 CC=g++
2 RANLIB=ranlib
3
4 LIBSRC=osm.cpp
5 LIBOBJ=$(LIBSRC:.cpp=.o)
6
7 INCS=-I.
8 CFLAGS = -Wall -g $(INCS)
9 LOADLIBES = -L./
10
11 OSMLIB = libosm.a
12 TARGETS = $(OSMLIB)
13
14 TAR=tar
15 TARFLAGS=-cvf
16 TARNAME=ex1.tar
17 TARSRC=$(LIBSRC) Makefile README
18
19 all: $(TARGETS)
20
21
22 $(TARGETS): $(LIBOBJ)
23     $(AR) $(ARFLAGS) $@ $^
24     $(RANLIB) $@
25
26 clean:
27     $(RM) $(TARGETS) $(OSMLIB) $(OBJ) $(LIBOBJ) *~ *core
28
29 depend:
30     makedepend -- $(CFLAGS) -- $(SRC) $(LIBSRC)
31
32 tar:
33     $(TAR) $(TARFLAGS) $(TARNAME) $(TARSRC)
```

3 osm.cpp

```
1  #include "osm.h"
2  #include <vector>
3  #include <sys/time.h>
4  #include <cstdlib>
5  #include <unistd.h>
6  #include <fcntl.h>
7
8  #define FINISH_SUCCESS 0
9  #define FINISH_ERROR -1
10 #define INVALID_ITER 0
11 #define DEFAULT_ITER 1000
12 #define MICRO_TO_NANO(x) ((x) * 1000)
13 #define SEC_TO_MICRO(x) ((x) * 1000000)
14 #define FILE_FLAGS O_CREAT | O_RDWR
15 #define FILE_NAME_1 "/tmp/.ransha"
16 #define FILE_NAME_2 "/tmp/.ransha1"
17 #define MAX_NAME_LEN 64
18
19 // g for global
20 static int gfd1 = 0, gfd2 = 0; // fd for file-descriptor
21 static char *gmachineName;
22
23 /* Initialization function that the user must call
24  * before running any other library function.
25  * The function may, for example, allocate memory or
26  * create/open files.
27  * Returns 0 upon success and -1 on failure
28  */
29 int osm_init()
30 {
31     // Attempt to open/create two files, if one fails the other isn't opened
32     // (logical or)
33     if (((gfd1 = open(FILE_NAME_1, FILE_FLAGS)) == FINISH_ERROR) ||
34         ((gfd2 = open(FILE_NAME_2, FILE_FLAGS)) == FINISH_ERROR))
35     {
36         return FINISH_ERROR;
37     }
38     // Allocate memory for hostname
39     gmachineName = new char[MAX_NAME_LEN];
40     return FINISH_SUCCESS;
41 }
42
43 /* finalizer function that the user must call
44  * after running any other library function.
45  * The function may, for example, free memory or
46  * close/delete files.
47  * Returns 0 upon success and -1 on failure
48  */
49 int osm_finalizer()
50 {
51     // Free allocated memory
52     delete gmachineName;
53
54     // Bitwise or to ensure both files are at least attempted to be closed
55     if ((close(gfd1) == FINISH_ERROR) | (close(gfd2) == FINISH_ERROR))
56     {
57         return FINISH_ERROR;
58     }
59     // Remove files.
```

```

60     if ((unlink(FILE_NAME_1) == FINISH_ERROR) |
61         (unlink(FILE_NAME_2) == FINISH_ERROR))
62     {
63         return FINISH_ERROR;
64     }
65     return FINISH_SUCCESS;
66 }
67
68 /*
69  * Empty functions for the function-call time measurements.
70  */
71 void emptyFunc() {}
72 void emptyFunc2() {}
73
74 /*
75  * Subtracts the second timeval struct from the first to get the difference
76  * between them in microseconds, then converts to nano-seconds.
77  */
78 double timeDiffInNano(timeval a, timeval b)
79 {
80     return MICRO_TO_NANO(SEC_TO_MICRO(a.tv_sec - b.tv_sec) \
81                          + a.tv_usec - b.tv_usec);
82 }
83
84 /* Time measurement function for a simple arithmetic operation.
85  * returns time in nano-seconds upon success,
86  * and -1 upon failure.
87  */
88 double osm_operation_time(unsigned int iterations)
89 {
90     // Ensure non-zero iterations
91     if (iterations == INVALID_ITER)
92     {
93         iterations = DEFAULT_ITER;
94     }
95
96     std::srand(0);
97     int a, b, c, d, j;
98     struct timeval st, et;
99     double tv, total = 0;
100
101     for (int i = 0; i < iterations; ++i)
102     {
103         // Get random numbers
104         a = std::rand();
105         b = std::rand();
106         c = std::rand();
107         d = std::rand();
108         if (gettimeofday(&st, NULL) != 0) // Start measure
109             return FINISH_ERROR;
110         // Do a lot of basic operations on random numbers */
111         for (j = 0; j < DEFAULT_ITER; ++j) {
112             //
113             a + b;
114             a & b;
115             a | b;
116             b + d;
117             b & d;
118             b | d;
119             a + c;
120             a & c;
121             a | c;
122             b + c;
123             b & c;
124             b | c;
125             a + d;
126             a & d;
127             a | d;

```

```

128     }
129     if (gettimeofday(&et, NULL) != 0) // End measure
130         return FINISH_ERROR;
131     // Calculate the time took (and convert to double)
132     tv = timeDiffInNano(et, st);
133     tv /= DEFAULT_ITER * 15; // 15 is the number of operations in
134                             // each inner loop iteration
135     total += tv; // accumulative measured time
136 }
137 return total / iterations; // Average time
138 }
139
140 /* Time measurement function for an empty function call.
141  * returns time in nano-seconds upon success,
142  * and -1 upon failure.
143  */
144 double osm_function_time(unsigned int iterations)
145 {
146     // Ensure non-zero iterations
147     if (iterations == INVALID_ITER)
148     {
149         iterations = DEFAULT_ITER;
150     }
151
152     int j;
153     struct timeval st, et;
154     double tv, total = 0;
155
156     for (int i = 0; i < iterations; ++i)
157     {
158         if (gettimeofday(&st, NULL) != 0)
159             return FINISH_ERROR;
160         for (j = 0; j < DEFAULT_ITER / 2; ++j)
161         {
162             emptyFunc();
163             emptyFunc2();
164             emptyFunc();
165             emptyFunc2();
166             emptyFunc();
167             emptyFunc2();
168             emptyFunc();
169             emptyFunc2();
170             emptyFunc();
171             emptyFunc2();
172             emptyFunc();
173             emptyFunc2();
174         }
175         if (gettimeofday(&et, NULL) != 0)
176             return FINISH_ERROR;
177         tv = timeDiffInNano(et, st);
178         tv /= DEFAULT_ITER / 2 * 12; // 12 for number of function calls in each
179                                     // iteration, DEFAULT_ITER / 2 iterations
180         total += tv;
181     }
182     return total / iterations;
183 }
184
185 /* Time measurement function for an empty trap into the operating system.
186  * returns time in nano-seconds upon success,
187  * and -1 upon failure.
188  */
189 double osm_syscall_time(unsigned int iterations)
190 {
191     // Ensure non-zero iterations
192     if (iterations == INVALID_ITER)
193     {
194         iterations = DEFAULT_ITER;
195     }

```

```

196
197     struct timeval st, et;
198     double tv, total = 0;
199
200     for (int i = 0; i < iterations; ++i)
201     {
202         if (gettimeofday(&st, NULL) != 0)
203             return FINISH_ERROR;
204         OSM_NULLSYSCALL;
205         OSM_NULLSYSCALL;
206         OSM_NULLSYSCALL;
207         OSM_NULLSYSCALL;
208         OSM_NULLSYSCALL;
209         if (gettimeofday(&et, NULL) != 0)
210             return FINISH_ERROR;
211         tv = timeDiffInNano(et, st);
212         tv /= 5; // 5 empty traps were made
213         total += tv;
214     }
215     return total / iterations;
216 }
217
218 /* Time measurement function for accessing the disk.
219 * returns time in nano-seconds upon success,
220 * and -1 upon failure.
221 */
222 double osm_disk_time(unsigned int iterations)
223 {
224     // Ensure non-zero iterations
225     if (iterations == INVALID_ITER)
226     {
227         iterations = DEFAULT_ITER;
228     }
229     int buffSize = 64;
230     struct timeval st, et;
231     double tv, total = 0;
232     char oneByte[buffSize];
233
234     for (int i = 0; i < iterations; ++i)
235     {
236         if (gettimeofday(&st, NULL) != 0)
237             return FINISH_ERROR;
238         // Try to write to both files
239         if ((write(gfd1, oneByte, 1) == FINISH_ERROR) ||
240             (write(gfd2, oneByte, 1) == FINISH_ERROR) ||
241             (read(gfd1, oneByte, 1) == FINISH_ERROR) ||
242             (read(gfd2, oneByte, 1) == FINISH_ERROR))
243         {
244             return FINISH_ERROR;
245         }
246         if (gettimeofday(&et, NULL) != 0)
247             return FINISH_ERROR;
248         tv = timeDiffInNano(et, st);
249         tv /= 4; // 4 is the number of "disk access"es performed
250         total += tv;
251     }
252     return total / iterations;
253 }
254
255 /*
256 * Measure all times and return a struct with the relevant data.
257 * Each field is set to -1 upon error or the correct value otherwise,
258 * except machineName which is set to the null-char '\0' upon failure.
259 */
260 timeMeasurmentStructure measureTimes(unsigned int operation_iterations,
261                                     unsigned int function_iterations,
262                                     unsigned int syscall_iterations,
263                                     unsigned int disk_iterations)

```



```

264 {
265     timeMeasurmentStructure retVal;
266     // Get host name
267     if (gethostname(gmachineName, MAX_NAME_LEN) != 0)
268     {
269         gmachineName = '\0';
270     }
271     retVal.machineName = gmachineName;
272
273     // Measure times
274     retVal.instructionTimeNanoSecond = \
275     osm_operation_time(operation_iterations);
276     retVal.functionTimeNanoSecond = osm_function_time(function_iterations);
277     retVal.trapTimeNanoSecond = osm_syscall_time(syscall_iterations);
278     retVal.diskTimeNanoSecond = osm_disk_time(disk_iterations);
279
280     // If basic op time is 0 for some reason, all ratios are
281     // invalid (divided by 0), therefore set to -1
282     if (retVal.instructionTimeNanoSecond == 0 ||
283         retVal.instructionTimeNanoSecond == FINISH_ERROR)
284     {
285         retVal.functionInstructionRatio = retVal.trapInstructionRatio = \
286         retVal.diskInstructionRatio = -1;
287     }
288     else
289     {
290         retVal.functionInstructionRatio = retVal.functionTimeNanoSecond / \
291         retVal.instructionTimeNanoSecond;
292         retVal.trapInstructionRatio = retVal.trapTimeNanoSecond / \
293         retVal.instructionTimeNanoSecond;
294         retVal.diskInstructionRatio = retVal.diskTimeNanoSecond / \
295         retVal.instructionTimeNanoSecond;
296     }
297     return retVal;
298 }

```